

# AI 2024

## Week 4 Task

### 1 Training an ANN

In this task, we will explore a simple Artificial Neural Network (ANN) with one hidden layer. The network has three input features, denoted as  $x_1$ ,  $x_2$ , and  $x_3$ , and one output  $\hat{y}$ . The architecture of the network is depicted in Figure 1.




Figure 1: A simple neural network with one hidden layer and three input features.

#### 1.1 Network Structure and Cost Function

1. **Input Layer:** Consists of three input neurons ( $x_1, x_2, x_3$ ).
2. **Hidden Layer:** Contains four neurons. Each neuron in the hidden layer calculates a weighted sum of all input features, adds a bias term, and applies a tanh activation function:

$$z_j^{(1)} = \sum_{i=1}^3 w_{ij}^{(1)} x_i + b_j^{(1)}, \quad a_j^{(1)} = \tanh(z_j^{(1)}) \quad \text{for } j = 1, 2, 3, 4.$$

3. **Output Layer:** Contains a single neuron that computes a weighted sum of all hidden layer outputs, adds a bias term, and applies a Sigmoid activation function:

$$z^{(2)} = \sum_{j=1}^4 w_j^{(2)} a_j^{(1)} + b^{(2)}, \quad \hat{y} = \sigma(z^{(2)}) = \frac{1}{1 + e^{-z^{(2)}}}.$$

4. **Cost Function:** The cost function  $J$  is defined as the average of the binary cross-entropy loss over all training examples. For a single training example, the binary cross-entropy loss  $\mathcal{L}$  is given by:

$$\mathcal{L} = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})). \quad (1)$$

**Notation:**

- $w_{ij}^{(l)}$ : Weight connecting neuron  $i$  in layer  $l - 1$  to neuron  $j$  in layer  $l$ .
- $b_j^{(l)}$ : Bias term for neuron  $j$  in layer  $l$ .
- $z_j^{(l)}$ : Weighted input to neuron  $j$  in layer  $l$  before activation.
- $a_j^{(l)}$ : Output of neuron  $j$  in layer  $l$  after applying the activation function.

**1.2 Initialization of Weights and Data Generation**

We will use random data points with three features and a corresponding label (0 or 1). The generated data points are:

| $\mathbf{x}$       | $y$ |
|--------------------|-----|
| (0.23, 0.45, 0.67) | 1   |
| (0.89, 0.12, 0.34) | 0   |
| (0.56, 0.78, 0.12) | 1   |
| (0.34, 0.88, 0.54) | 0   |
| (0.98, 0.76, 0.65) | 1   |

We initialize the weights and biases randomly as follows:

$$\mathbf{w}^{(1)} = \begin{bmatrix} 0.45 & -0.12 & 0.78 \\ 0.05 & 0.35 & -0.22 \\ -0.55 & 0.11 & 0.67 \\ 0.72 & -0.85 & 0.45 \end{bmatrix}, \quad \mathbf{b}^{(1)} = \begin{bmatrix} 0.1 \\ -0.1 \\ 0.2 \\ -0.2 \end{bmatrix},$$

$$\mathbf{w}^{(2)} = \begin{bmatrix} 0.25 \\ -0.35 \\ 0.15 \\ -0.05 \end{bmatrix}, \quad \mathbf{b}^{(2)} = 0.3.$$

**1.3 Tasks****1.3.1 Gradients and Partial Derivatives**

Calculate the partial derivatives required for backpropagation:

1. Derivative of the loss  $\mathcal{L}$  with respect to the output of the network  $\hat{y}$ :

$$\frac{\partial \mathcal{L}}{\partial \hat{y}}$$

2. Derivative of the output activation (Sigmoid) with respect to its input:

$$\frac{\partial \hat{y}}{\partial z^{(2)}}$$

3. Derivative of the activation of hidden layer neurons with respect to their weighted input (tanh):

$$\frac{\partial a_j^{(1)}}{\partial z_j^{(1)}}$$

4. Derivative of the cost function  $J$  with respect to the output neuron input  $z^{(2)}$  using the loss function from Equation (1):

$$\frac{\partial J}{\partial z^{(2)}}$$

5. Derivative of the loss with respect to the weights of the hidden layer:

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^{(1)}}$$

6. Derivative of the loss with respect to the bias terms in the hidden layer:

$$\frac{\partial \mathcal{L}}{\partial b_j^{(1)}}$$

7. Derivative of the loss with respect to the weights of the output layer:

$$\frac{\partial \mathcal{L}}{\partial w_j^{(2)}}$$

8. Derivative of the loss with respect to the bias term in the output layer:

$$\frac{\partial \mathcal{L}}{\partial b^{(2)}}$$

### 1.3.2 Forward Pass

Perform a forward pass through the network using the first data point and the weights and biases provided in the Initialization of Weights and Data Generation subsection:

1. Calculate the values  $z_j^{(1)}$  for each neuron in the hidden layer.
2. Apply the tanh activation function to compute  $a_j^{(1)}$ .
3. Compute the output neuron input  $z^{(2)}$ .
4. Apply the Sigmoid activation function to compute  $\hat{y}$ .

### 1.3.3 Backpropagation and Training

Using the backpropagation algorithm:

1. Compute the gradients for each weight and bias using the partial derivatives and the loss function from Equation (1).
2. Update the weights and biases using Gradient Descent:

$$w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - \alpha \frac{\partial \mathcal{L}}{\partial w_{ij}^{(l)}}, \quad b_j^{(l)} \leftarrow b_j^{(l)} - \alpha \frac{\partial \mathcal{L}}{\partial b_j^{(l)}}$$

where  $\alpha = 0.05$  is the learning rate.

## 2 Implementing an Artificial Neural Network for the Titanic Dataset

In this section, you will implement a simple Artificial Neural Network (ANN) from scratch using NumPy, without using any deep learning frameworks such as TensorFlow or PyTorch. The objective is to train a model that predicts whether a passenger survived the Titanic disaster based on their characteristics. The network will have a maximum of two layers (one hidden layer and one output layer).

### 2.1 Steps to Implement the ANN

Follow these steps to complete the task:

#### 2.1.1 Download the Datasets

Download the train and test datasets from the Titanic competition on Kaggle.

#### 2.1.2 Load the Datasets

Load the downloaded datasets into your Python environment using libraries such as pandas.

#### 2.1.3 Perform Exploratory Data Analysis (EDA)

Conduct exploratory data analysis (EDA) on the datasets to understand the features and their distributions. In particular, focus on:

- Identifying and handling missing values.
- Normalizing numerical features to ensure they are on a similar scale.
- Encoding categorical features into numerical values if necessary.

You can refer to your Week 1 task results for insights and methods on EDA.

#### 2.1.4 Divide the Train Set into Feature Matrix and Labels

Divide the training dataset into two separate arrays:

- $\mathbf{X}_{\text{train}}$ : A matrix containing all the input features for the training examples.
- $\mathbf{Y}_{\text{train}}$ : A vector containing the corresponding labels (0 for not survived, 1 for survived).

Similarly, prepare the test dataset features as  $\mathbf{X}_{\text{test}}$ .

#### 2.1.5 Initialize Weights and Biases

Randomly initialize the weights and biases for each layer in your neural network. Use small random values to start the optimization process.

#### 2.1.6 Implement Forward Pass Function

Write a function to perform a forward pass through the network. This function should compute the outputs (activations) for each layer given an input feature vector.

#### 2.1.7 Implement Backpropagation Function

Write a function to perform backpropagation. This function should compute the gradients (partial derivatives) of the loss function with respect to each weight and bias.

#### 2.1.8 Update Parameters Function

Write a function to update the network's parameters (weights and biases) using the gradients computed in the backpropagation step. The update should use gradient descent and a specified learning rate.

#### 2.1.9 Prediction Function

Write a function to predict if a passenger has survived or not. The function should take an input feature vector, perform a forward pass, and output a binary prediction (0 or 1).

#### 2.1.10 Training and Evaluation Function

Combine the above functions into a single function that:

- Takes  $\mathbf{X}_{\text{train}}$ ,  $\mathbf{Y}_{\text{train}}$ ,  $\mathbf{X}_{\text{test}}$ , `num_iterations`, and `learning_rate` as inputs.
- Trains the model over the specified number of iterations.
- Logs the loss on the training set for each iteration.
- Saves the predictions for the test set in a CSV file formatted according to the Kaggle submission requirements.

**Note:** Ensure that your code is vectorized and utilizes NumPy for all matrix and vector operations to optimize performance.

#### **2.1.11 Upload Results to Kaggle**

Upload your prediction CSV file to Kaggle to evaluate your model's accuracy. Record the score for reporting purposes.

#### **2.1.12 Upload to IUTBox**

Finally, upload your code, the Kaggle predictions CSV file, and the accuracy score obtained from Kaggle to IUTBox as instructed.

### **2.2 Optional Steps**

- Experiment with different architectures (e.g., changing the number of neurons in the hidden layer) to see their impact on performance.
- Plot the loss curve to visualize how the training progresses over time.
- Test different learning rates and observe their effects on model convergence.